

Bayesian Missing Data - Multiple Imputation Methods

Presented to



California State University, Fullerton

Math 538 Fall 2023

Dr. Poynor

Prepared by

Paul LOPEZ

Emilio VASQUEZ

December 07, 2023

1 Instructions

blah blah blah

1.1 Introduction

Introduce your topic: Provide the background and/or foundation of your topic (e.g. the type of data, application, or setting, etc). Define notation. You may have a broader perspective here.

Missing data is an unavoidable issue that arises in many real-world datasets across various fields, including finance, healthcare, and social sciences. When data are missing, it can introduce biases and lead to invalid statistical inferences if the missing values are simply ignored or eliminated. Multiple imputation is a principled approach for handling missing data that involves creating multiple complete versions of the incomplete dataset, with the missing values imputed through simulated draws from an appropriate model. The key insight is that missing data uncertainty can be represented by generating multiple imputed datasets.

In the Bayesian approach to multiple imputation, prior distributions are specified for the model parameters and missing data, representing our beliefs before seeing the data. The posterior distributions of the parameters and missing values are estimated by fitting the model to the observed data using Markov chain Monte Carlo (MCMC) methods. This allows simulating multiple imputed datasets by drawing missing values from their posterior predictive distributions. The completed datasets can then be analyzed using standard complete-data methods, with final estimates pooled across the imputed datasets to incorporate missing data uncertainty. Bayesian multiple imputation provides a flexible framework for handling missing data while properly representing imputation uncertainty.

A key advantage of the Bayesian approach to multiple imputation is that it allows incorporating appropriate prior information and modeling complex relationships between variables. For example, hierarchical priors can be used to share information between related parameters or models, improving estimates for variables with limited data. Bayesian models like mixtures and nonparametric models provide flexibility to adapt to complex patterns in the data. MCMC provides a convenient computational approach for fitting Bayesian models with missing data, avoiding analytical intractability. The posterior predictive distribution for the missing values conditions on both the observed data and model parameters, ensuring proper imputation uncertainty. Practical implementations utilize packages like `mice` in R and `mi` in Python, which automate iterative Bayesian modeling, imputation, and analysis. Overall, Bayesian multiple imputation provides a robust approach for handling missing data that accounts for uncertainty and allows incorporating flexible modeling and prior information about the data structure.

1.2 Motivation

Motivation/Contribution: Motivate the importance or/and use of the specific model or process you are presenting. Note that topics may have several motivations and contributions, due to time

constraints perhaps focus on one or two.

Motivation 1: Missing data is very common in financial datasets, especially due to non-response on surveys or forms. For example, customers may not fill out all the fields on a loan application or investors may skip questions on an investment risk tolerance questionnaire. Simply throwing out incomplete records would result in losing valuable information and potential biases. Multiple imputation allows you to predict missing values based on patterns in the observed data. This provides a principled way to fill in missing values while accounting for uncertainty, rather than ad hoc methods like mean imputation. Some models that can be used are multivariate normal imputation or chained equations/sequential regression approaches.

Motivation 2: Many financial models require complete data, such as credit risk models that use various borrower characteristics as predictors. Multiple imputation can be used to impute missing predictor values so that all records can be used, improving model accuracy and avoiding selection bias. Specific models include imputation based on Bayesian networks, which encode conditional dependencies between variables. Tree-based methods like random forests can also be used to model complex nonlinear relationships for imputation. The multiple imputed datasets can then be used to train the financial models, with results appropriately combined across the imputed datasets. This allows building more robust models on messy real-world data.

Motivation 3: In portfolio optimization and risk models, it is important to have complete data on asset returns. However, there may be gaps in time series data due to illiquidity or data errors. Not accounting for these gaps can bias risk measures like volatility and correlations. Multiple imputation can be used to fill in missing return data in a way that represents the uncertainty. Common approaches are based on multivariate normal models or autoregressive time series models. The imputed return datasets can then be used as input to optimize portfolios or estimate covariance matrices for risk management. This allows building more accurate financial models and avoiding potential pitfalls from ignoring missing data. Multiple imputation provides a flexible framework to handle missing values while properly quantifying the uncertainty in the imputed values.

1.3 Structure of Model/Process

Provide the structure of the Model/Process: Present the structure of the Model/Process. Perhaps comment on the role of the parameters. What prior distribution(s) are you going to use.

The model structure is a standard Bayesian regression model relating a continuous response Y to predictor variables X :

$$Y \sim N(\mu, \sigma)$$

$$\mu = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots$$

where:

- Y is the response variable with some values missing
- X_1, X_2 , etc are the predictor variables
- β_0 is the intercept
- β_1, β_2 , etc are the regression coefficients
- σ is the error standard deviation

In addition, we have missingness indicator variables R for each observation, where R = 1 if Y is observed and R = 0 if Y is missing for that observation.

The priors on the parameters $\beta_0, \beta_1, \beta_2$, and σ are weakly informative normals or half-normals:

$$\beta_0 \sim N(0, 10)$$

$$\beta_1 \sim N(0, 2)$$

$$\beta_2 \sim N(0, 2)$$

$$\sigma \sim HalfNormal(0, 10)$$

For the missing Y values, we treat them as parameters to be estimated, with priors based on the observed data likelihood:

$$Y_{miss} \sim N(\mu, \sigma)$$

where μ and σ come from fitting the model on the observed data. This makes the missing values exchangeable with the observed data.

Multiple imputation is done by drawing missing Y values from their posterior predictive distribution based on the fitted model. Multiple datasets are created by repeating this process and used to account for missing data uncertainty.

In summary, the model leverages Bayesian regression, weakly informative priors, and treats missing values as parameters to plausibly fill in gaps while quantifying uncertainty. The multiple imputed datasets integrate over the posterior distribution of the missing data.

1.4 Example

Example: You may use a data example, simulation, or examine properties simply by simulating from the Model/Process itself. This is a Bayesian class, so you should discuss the choice of prior values. Discuss results of Example. Note you may use built functions in your example, but please provide the details in your presentation/report as to what the function is doing.

The model being fit in this analysis is a Bayesian linear regression model with time-series data. The formula for the model is:

$$return_t = \beta_0 + \beta_1 laggedreturn_{t-1} + \epsilon_t$$

- $return_t$ is the dependent variable, representing the log return of the AAPL stock on day t. It is calculated as the difference in the logarithm of the closing prices between two consecutive days, i.e., $\log(close_t) - \log(close_{t-1})$.
- $laggedreturn_{t-1}$ is the independent variable, which is the log return of the AAPL stock on the previous day t-1.
- β_0 is the intercept of the regression line, which represents the expected value of $return_t$, when $laggedreturn_{t-1}$ is zero.
- β_1 is the slope coefficient, indicating how much $return_t$ is expected to increase when $laggedreturn_{t-1}$ increases by one unit.
- ϵ_t is the error term, which accounts for the variability in $return_t$ that is not explained by $laggedreturn_{t-1}$.

The Bayesian aspect of this model comes from the use of priors and the Bayesian inference process. Instead of just finding point estimates for β_0 and β_1 as in classical regression, the Bayesian approach estimates the entire posterior distributions for these parameters based on the prior distributions and the observed data.

Below are the steps taken for this analysis:

1. Data Acquisition: The `tq_get` function from the `tidyquant` package retrieves historical stock price data for AAPL from an online source.

2. Data Preparation for Amputation:

- A subset of the retrieved data is selected, focusing on the ‘close’ (closing price) and ‘volume’ (number of shares traded) columns.
- The `set.seed` function sets the random number generator’s seed to ensure reproducibility of the results.

3. Data Amputation:

The `mice::ampute` function artificially introduces missing values into the dataset to simulate incomplete data, which is common in real-world scenarios. The `prop` argument specifies the proportion of data to be made missing, and the `mech` argument specifies the missingness mechanism as ‘MAR’ (Missing At Random).

4. Data Preparation Post-Amputation:

- The stocks dataframe is updated with the amputed ‘close’ prices.
- New variables ‘return’ and ‘lagged_return’ are calculated using the `dplyr` package. These represent the daily returns of the stock and the previous day’s return, respectively. The `na.omit` function removes any rows with missing values which can’t be used in the subsequent regression analysis.

5. Data Imputation:

The `mice` function is used to impute missing values in the dataset. It uses Predictive Mean Matching (‘pmm’) method to fill in missing ‘close’ prices. Multiple imputed datasets are generated (specified by `m = 5`) to account for the uncertainty of the imputation process.

The PMM Process for this `mice` (Multivariate Imputation by Chained Equations) function works as follows:

- i) Fitting of the initial model: We have our variables generically as X (predictors) and Y (variable with missing values). We fit a regression model:

$$\hat{Y} = f(X; \theta), \text{ where } \theta \text{ are the estimated parameters}$$

- ii) For a missing Y_i , we find the set of cases (called donors in the package) D_i where $\hat{Y}_d \in D_i$ are closest to \hat{Y}_i .
- iii) We now randomly select a Y_d from D_i to impute Y_i . We have the option to do a straight imputation, or we can choose n of them and take an average. In our case we did straight imputation, and the selection from this set D gives each candidate an equal chance of being selected. From what we could gather from the documentation this probability process was akin to a uniform probability distribution over the pool of potential donors.

This PMM method is a standalone, non-Bayesian imputation method. The imputed datasets generated are then used further using Bayesian methods in this project. By creating multiple imputations ($m = 5$), **mice** acknowledges and represents the uncertainty inherent in the imputation process.

6. Bayesian Regression Model:

- Priors are defined for the Bayesian regression model using the **prior** function. The priors reflect our beliefs about the parameters before seeing the data. In this case, normal priors are set for the intercept and slope (return and lagged_return coefficients).
- The **brm** function from the **brms** package fits a Bayesian regression model to each imputed dataset. The models predict 'return' based on 'lagged_return'. The **iter**, **warmup**, and **chains** arguments control the Markov Chain Monte Carlo (MCMC) sampling process used to estimate the posterior distribution of the model parameters.
- Intercept Prior: A normal distribution with a mean of 0 and a standard deviation of 2. This is specified by `prior(normal(0, 2), class = "Intercept")`. It indicates that before looking at the data, you believe the average daily return when the lagged return is zero is likely to be around 0, but there is some uncertainty, which is captured by the standard deviation of 2.
- Slope Prior: A normal distribution with a mean of 0 and a standard deviation of 1. This is specified by `prior(normal(0, 1), class = "b")`. It suggests that before analyzing the data, you expect the impact of the previous day's return on today's return to be small, as indicated by the mean of 0. The standard deviation of 1 reflects your uncertainty about this expectation.

The posterior distribution combines the prior distribution with the likelihood of the observed data to update our beliefs about the model parameters. After fitting the model using **brm**, which internally uses Hamiltonian Monte Carlo (HMC) sampling, a type of MCMC, you obtain a distribution for each parameter that reflects all the information from the priors and the data.

7. Combining of Model Results:

- Parameter estimates and standard errors are extracted from each fitted model. These estimates are combined by calculating the mean estimate and pooled standard error for each parameter across all imputed datasets.

8. Summary and Analysis:

The combined results are printed out for interpretation.

This analysis allows for Bayesian inference on time-series data with missing values. The choice of prior values should be justified based on domain knowledge or previous research. Since normal priors are used for both the intercept and slope, it implies a belief that the coefficients are likely

to be around 0 but allowing for some variation.

The `mice::ampute` function is crucial because it enables researchers to understand how well their imputation and analysis methods work when some data is missing, which is a common occurrence in real-world datasets.

The `brm` function is central to the Bayesian approach, as it fits a model within the Bayesian framework using Hamiltonian Monte Carlo, a type of MCMC sampling. The function returns a wealth of information about the model, including estimates of the posterior distribution of the model parameters, which reflects both the data and the priors.

```
# Get Apple stock data
stocks <- tq_get("AAPL", get = "stock.prices", from = "2013-01-01")

# Select multiple columns for the amputation process
stocks_for_ampute <- stocks[, c("close", "volume")]

# Introduce missing values
set.seed(555)
amputed_data <- mice::ampute(stocks_for_ampute, prop = 0.3, mech = "MAR")

# Update the stocks with the amputed data for 'close'
stocks$close <- amputed_data$amp[, "close"]

# Ensure all necessary columns are present after amputation
stocks <- as.data.frame(stocks)

# Calculate 'return' and 'lagged_return' for your model
stocks <- stocks %>%
  mutate(
    return = c(NA, diff(log(close))),
    lagged_return = lag(return, 1)
  ) %>%
  na.omit() # Remove rows with NAs

# Impute missing values using Predictive Mean Matching
imputed <- mice(data = stocks, m = 5, maxit = 20, method = 'pmm', seed = 123)

##
## iter imp variable
```

##	1	1
##	1	2
##	1	3
##	1	4
##	1	5
##	2	1
##	2	2
##	2	3
##	2	4
##	2	5
##	3	1
##	3	2
##	3	3
##	3	4
##	3	5
##	4	1
##	4	2
##	4	3
##	4	4
##	4	5
##	5	1
##	5	2
##	5	3
##	5	4
##	5	5
##	6	1
##	6	2
##	6	3
##	6	4
##	6	5
##	7	1
##	7	2
##	7	3
##	7	4
##	7	5
##	8	1
##	8	2

##	8	3
##	8	4
##	8	5
##	9	1
##	9	2
##	9	3
##	9	4
##	9	5
##	10	1
##	10	2
##	10	3
##	10	4
##	10	5
##	11	1
##	11	2
##	11	3
##	11	4
##	11	5
##	12	1
##	12	2
##	12	3
##	12	4
##	12	5
##	13	1
##	13	2
##	13	3
##	13	4
##	13	5
##	14	1
##	14	2
##	14	3
##	14	4
##	14	5
##	15	1
##	15	2
##	15	3
##	15	4

```
## 15 5
## 16 1
## 16 2
## 16 3
## 16 4
## 16 5
## 17 1
## 17 2
## 17 3
## 17 4
## 17 5
## 18 1
## 18 2
## 18 3
## 18 4
## 18 5
## 19 1
## 19 2
## 19 3
## 19 4
## 19 5
## 20 1
## 20 2
## 20 3
## 20 4
## 20 5
```

```
## Warning: Number of logged events: 5
```

```
# Extract imputed datasets
```

```
imp_datasets <- lapply(1:5, function(i) complete(imputed, i))
```

```
# Define the priors for the Bayesian regression model
```

```
priors <- c(
  prior(normal(0, 2), class = "Intercept"),
  prior(normal(0, 1), class = "b")
)
```

```
# Initialize an empty list to store the models
```

```
fits <- vector("list", length = 5)

# Fit a Bayesian regression model to each imputed dataset and check for errors
for (i in seq_along(imp_datasets)) {
  dataset <- imp_datasets[[i]]
  fit <- tryCatch(
    {
      model <- brm(return ~ lagged_return, data = dataset,
                   iter = 1000, warmup = 500, chains = 2, prior = priors)
      model # Return the model
    },
    error = function(e) {
      cat(sprintf("Error in fitting model %d: %s\n", i, e$message))
      NULL # Return NULL if there was an error
    }
  )
  fits[[i]] <- fit
}
```

```
## Compiling Stan program...
```

```
## Start sampling
```

```
##
```

```
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 1).
```

```
## Chain 1:
```

```
## Chain 1: Gradient evaluation took 4e-05 seconds
```

```
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.4 seconds
```

```
## Chain 1: Adjust your expectations accordingly!
```

```
## Chain 1:
```

```
## Chain 1:
```

```
## Chain 1: Iteration: 1 / 1000 [ 0%] (Warmup)
```

```
## Chain 1: Iteration: 100 / 1000 [ 10%] (Warmup)
```

```
## Chain 1: Iteration: 200 / 1000 [ 20%] (Warmup)
```

```
## Chain 1: Iteration: 300 / 1000 [ 30%] (Warmup)
```

```
## Chain 1: Iteration: 400 / 1000 [ 40%] (Warmup)
```

```
## Chain 1: Iteration: 500 / 1000 [ 50%] (Warmup)
```

```
## Chain 1: Iteration: 501 / 1000 [ 50%] (Sampling)
```

```
## Chain 1: Iteration: 600 / 1000 [ 60%] (Sampling)
```

```

## Chain 1: Iteration: 700 / 1000 [ 70%] (Sampling)
## Chain 1: Iteration: 800 / 1000 [ 80%] (Sampling)
## Chain 1: Iteration: 900 / 1000 [ 90%] (Sampling)
## Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.129 seconds (Warm-up)
## Chain 1: 0.104 seconds (Sampling)
## Chain 1: 0.233 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 8e-06 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.08 seconds
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration: 1 / 1000 [ 0%] (Warmup)
## Chain 2: Iteration: 100 / 1000 [ 10%] (Warmup)
## Chain 2: Iteration: 200 / 1000 [ 20%] (Warmup)
## Chain 2: Iteration: 300 / 1000 [ 30%] (Warmup)
## Chain 2: Iteration: 400 / 1000 [ 40%] (Warmup)
## Chain 2: Iteration: 500 / 1000 [ 50%] (Warmup)
## Chain 2: Iteration: 501 / 1000 [ 50%] (Sampling)
## Chain 2: Iteration: 600 / 1000 [ 60%] (Sampling)
## Chain 2: Iteration: 700 / 1000 [ 70%] (Sampling)
## Chain 2: Iteration: 800 / 1000 [ 80%] (Sampling)
## Chain 2: Iteration: 900 / 1000 [ 90%] (Sampling)
## Chain 2: Iteration: 1000 / 1000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.145 seconds (Warm-up)
## Chain 2: 0.111 seconds (Sampling)
## Chain 2: 0.256 seconds (Total)
## Chain 2:
## Compiling Stan program...
## Start sampling
##

```

```
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 4.3e-05 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.43 seconds
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:    1 / 1000 [  0%] (Warmup)
## Chain 1: Iteration: 100 / 1000 [ 10%] (Warmup)
## Chain 1: Iteration: 200 / 1000 [ 20%] (Warmup)
## Chain 1: Iteration: 300 / 1000 [ 30%] (Warmup)
## Chain 1: Iteration: 400 / 1000 [ 40%] (Warmup)
## Chain 1: Iteration: 500 / 1000 [ 50%] (Warmup)
## Chain 1: Iteration: 501 / 1000 [ 50%] (Sampling)
## Chain 1: Iteration: 600 / 1000 [ 60%] (Sampling)
## Chain 1: Iteration: 700 / 1000 [ 70%] (Sampling)
## Chain 1: Iteration: 800 / 1000 [ 80%] (Sampling)
## Chain 1: Iteration: 900 / 1000 [ 90%] (Sampling)
## Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.116 seconds (Warm-up)
## Chain 1:                0.106 seconds (Sampling)
## Chain 1:                0.222 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 1.8e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.18 seconds
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration:    1 / 1000 [  0%] (Warmup)
## Chain 2: Iteration: 100 / 1000 [ 10%] (Warmup)
## Chain 2: Iteration: 200 / 1000 [ 20%] (Warmup)
## Chain 2: Iteration: 300 / 1000 [ 30%] (Warmup)
## Chain 2: Iteration: 400 / 1000 [ 40%] (Warmup)
```

```
## Chain 2: Iteration: 500 / 1000 [ 50%] (Warmup)
## Chain 2: Iteration: 501 / 1000 [ 50%] (Sampling)
## Chain 2: Iteration: 600 / 1000 [ 60%] (Sampling)
## Chain 2: Iteration: 700 / 1000 [ 70%] (Sampling)
## Chain 2: Iteration: 800 / 1000 [ 80%] (Sampling)
## Chain 2: Iteration: 900 / 1000 [ 90%] (Sampling)
## Chain 2: Iteration: 1000 / 1000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.113 seconds (Warm-up)
## Chain 2: 0.121 seconds (Sampling)
## Chain 2: 0.234 seconds (Total)
## Chain 2:

## Compiling Stan program...
## Start sampling

##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 4.8e-05 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.48 seconds
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration: 1 / 1000 [ 0%] (Warmup)
## Chain 1: Iteration: 100 / 1000 [ 10%] (Warmup)
## Chain 1: Iteration: 200 / 1000 [ 20%] (Warmup)
## Chain 1: Iteration: 300 / 1000 [ 30%] (Warmup)
## Chain 1: Iteration: 400 / 1000 [ 40%] (Warmup)
## Chain 1: Iteration: 500 / 1000 [ 50%] (Warmup)
## Chain 1: Iteration: 501 / 1000 [ 50%] (Sampling)
## Chain 1: Iteration: 600 / 1000 [ 60%] (Sampling)
## Chain 1: Iteration: 700 / 1000 [ 70%] (Sampling)
## Chain 1: Iteration: 800 / 1000 [ 80%] (Sampling)
## Chain 1: Iteration: 900 / 1000 [ 90%] (Sampling)
## Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.068 seconds (Warm-up)
## Chain 1: 0.113 seconds (Sampling)
```



```

## Chain 1:                0.181 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 4.3e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.43 seconds
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration:    1 / 1000 [  0%] (Warmup)
## Chain 2: Iteration:  100 / 1000 [ 10%] (Warmup)
## Chain 2: Iteration:  200 / 1000 [ 20%] (Warmup)
## Chain 2: Iteration:  300 / 1000 [ 30%] (Warmup)
## Chain 2: Iteration:  400 / 1000 [ 40%] (Warmup)
## Chain 2: Iteration:  500 / 1000 [ 50%] (Warmup)
## Chain 2: Iteration:  501 / 1000 [ 50%] (Sampling)
## Chain 2: Iteration:  600 / 1000 [ 60%] (Sampling)
## Chain 2: Iteration:  700 / 1000 [ 70%] (Sampling)
## Chain 2: Iteration:  800 / 1000 [ 80%] (Sampling)
## Chain 2: Iteration:  900 / 1000 [ 90%] (Sampling)
## Chain 2: Iteration: 1000 / 1000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.079 seconds (Warm-up)
## Chain 2:                0.086 seconds (Sampling)
## Chain 2:                0.165 seconds (Total)
## Chain 2:

## Compiling Stan program...
## Start sampling

##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 6.5e-05 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.65 seconds
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:

```

```
## Chain 1: Iteration: 1 / 1000 [ 0%] (Warmup)
## Chain 1: Iteration: 100 / 1000 [ 10%] (Warmup)
## Chain 1: Iteration: 200 / 1000 [ 20%] (Warmup)
## Chain 1: Iteration: 300 / 1000 [ 30%] (Warmup)
## Chain 1: Iteration: 400 / 1000 [ 40%] (Warmup)
## Chain 1: Iteration: 500 / 1000 [ 50%] (Warmup)
## Chain 1: Iteration: 501 / 1000 [ 50%] (Sampling)
## Chain 1: Iteration: 600 / 1000 [ 60%] (Sampling)
## Chain 1: Iteration: 700 / 1000 [ 70%] (Sampling)
## Chain 1: Iteration: 800 / 1000 [ 80%] (Sampling)
## Chain 1: Iteration: 900 / 1000 [ 90%] (Sampling)
## Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.111 seconds (Warm-up)
## Chain 1: 0.102 seconds (Sampling)
## Chain 1: 0.213 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 1.5e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.15 seconds
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration: 1 / 1000 [ 0%] (Warmup)
## Chain 2: Iteration: 100 / 1000 [ 10%] (Warmup)
## Chain 2: Iteration: 200 / 1000 [ 20%] (Warmup)
## Chain 2: Iteration: 300 / 1000 [ 30%] (Warmup)
## Chain 2: Iteration: 400 / 1000 [ 40%] (Warmup)
## Chain 2: Iteration: 500 / 1000 [ 50%] (Warmup)
## Chain 2: Iteration: 501 / 1000 [ 50%] (Sampling)
## Chain 2: Iteration: 600 / 1000 [ 60%] (Sampling)
## Chain 2: Iteration: 700 / 1000 [ 70%] (Sampling)
## Chain 2: Iteration: 800 / 1000 [ 80%] (Sampling)
## Chain 2: Iteration: 900 / 1000 [ 90%] (Sampling)
## Chain 2: Iteration: 1000 / 1000 [100%] (Sampling)
```

```

## Chain 2:
## Chain 2: Elapsed Time: 0.1 seconds (Warm-up)
## Chain 2:          0.123 seconds (Sampling)
## Chain 2:          0.223 seconds (Total)
## Chain 2:

## Compiling Stan program...
## Start sampling

##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 1).
## Chain 1:
## Chain 1: Gradient evaluation took 5.3e-05 seconds
## Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.53 seconds
## Chain 1: Adjust your expectations accordingly!
## Chain 1:
## Chain 1:
## Chain 1: Iteration:   1 / 1000 [  0%] (Warmup)
## Chain 1: Iteration: 100 / 1000 [ 10%] (Warmup)
## Chain 1: Iteration: 200 / 1000 [ 20%] (Warmup)
## Chain 1: Iteration: 300 / 1000 [ 30%] (Warmup)
## Chain 1: Iteration: 400 / 1000 [ 40%] (Warmup)
## Chain 1: Iteration: 500 / 1000 [ 50%] (Warmup)
## Chain 1: Iteration: 501 / 1000 [ 50%] (Sampling)
## Chain 1: Iteration: 600 / 1000 [ 60%] (Sampling)
## Chain 1: Iteration: 700 / 1000 [ 70%] (Sampling)
## Chain 1: Iteration: 800 / 1000 [ 80%] (Sampling)
## Chain 1: Iteration: 900 / 1000 [ 90%] (Sampling)
## Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)
## Chain 1:
## Chain 1: Elapsed Time: 0.13 seconds (Warm-up)
## Chain 1:          0.149 seconds (Sampling)
## Chain 1:          0.279 seconds (Total)
## Chain 1:
##
## SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).
## Chain 2:
## Chain 2: Gradient evaluation took 4.3e-05 seconds
## Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.43 seconds

```

```
## Chain 2: Adjust your expectations accordingly!
## Chain 2:
## Chain 2:
## Chain 2: Iteration: 1 / 1000 [ 0%] (Warmup)
## Chain 2: Iteration: 100 / 1000 [ 10%] (Warmup)
## Chain 2: Iteration: 200 / 1000 [ 20%] (Warmup)
## Chain 2: Iteration: 300 / 1000 [ 30%] (Warmup)
## Chain 2: Iteration: 400 / 1000 [ 40%] (Warmup)
## Chain 2: Iteration: 500 / 1000 [ 50%] (Warmup)
## Chain 2: Iteration: 501 / 1000 [ 50%] (Sampling)
## Chain 2: Iteration: 600 / 1000 [ 60%] (Sampling)
## Chain 2: Iteration: 700 / 1000 [ 70%] (Sampling)
## Chain 2: Iteration: 800 / 1000 [ 80%] (Sampling)
## Chain 2: Iteration: 900 / 1000 [ 90%] (Sampling)
## Chain 2: Iteration: 1000 / 1000 [100%] (Sampling)
## Chain 2:
## Chain 2: Elapsed Time: 0.12 seconds (Warm-up)
## Chain 2: 0.148 seconds (Sampling)
## Chain 2: 0.268 seconds (Total)
## Chain 2:
```

```
# Initialize lists to store parameter estimates and standard errors
param_estimates <- list()
param_se <- list()

# Extract parameter estimates and standard errors from each model
for (i in seq_along(fits)) {
  if (!is.null(fits[[i]])) {
    # Extracting estimates
    est <- as.data.frame(summary(fits[[i]])$fixed)[, "Estimate"]
    param_estimates[[i]] <- est

    # Extracting standard errors
    se <- as.data.frame(summary(fits[[i]])$fixed)[, "Est.Error"]
    param_se[[i]] <- se
  }
}
```

```

# Calculate the mean of the parameter estimates
combined_estimates <- Reduce("+", param_estimates) / length(fits)

# Calculate the pooled standard error if necessary
combined_se <- sqrt(Reduce("+", lapply(param_se, `^`, 2))) / length(fits)

# Combine into a data frame
combined_results <- data.frame(Estimate = combined_estimates, Std.Error = combined_se)

# Print combined results
#print(combined_results)

# Initialize lists to store CIs
ci_lower <- list()
ci_upper <- list()

for (i in seq_along(fits)) {

  if (!is.null(fits[[i]])) {

    # Extract CIs
    ci <- as.data.frame(summary(fits[[i]])$fixed)[, c("l-95% CI", "u-95% CI")]

    ci_lower[[i]] <- ci[,1]
    ci_upper[[i]] <- ci[,2]

  }

}

# Combine CIs

ci_lower_combined <- Reduce("+", ci_lower) / length(fits)
ci_upper_combined <- Reduce("+", ci_upper) / length(fits)

combined_results$ci_lower <- ci_lower_combined
combined_results$ci_upper <- ci_upper_combined

```

```
print(combined_results)
```

```
##      Estimate   Std.Error    ci_lower    ci_upper
## 1  0.001010482 0.000175254  0.0002387094 0.001772605
## 2 -0.053271826 0.010511568 -0.0988787790 -0.006764432
```

```
# PLOT CODE WORKS BUT THE PLOTS ARE CRAP FOR NOW
```

```
# Create a sequence of x values for the density plots
```

```
x_intercept <- seq(min(unlist(param_estimates[[1]])) - 3 * max(unlist(param_se[[1]])),
                  max(unlist(param_estimates[[1]])) + 3 * max(unlist(param_se[[1]])), length.out = 100)
x_slope <- seq(min(unlist(param_estimates[[2]])) - 3 * max(unlist(param_se[[2]])),
              max(unlist(param_estimates[[2]])) + 3 * max(unlist(param_se[[2]])), length.out = 100)
```

```
# Function to generate the density data
```

```
generate_density_data <- function(estimates, ses, x_values) {
  densities <- mapply(function(estimate, se) {
    dnorm(x_values, mean = estimate, sd = se)
  }, estimate = estimates, se = ses)

  mean_densities <- colMeans(densities)
  data.frame(x = x_values, y = mean_densities)
}
```

```
# Generate the density data
```

```
density_data_intercept <- generate_density_data(unlist(param_estimates[[1]]), unlist(param_se[[1]]), x_intercept)
density_data_slope <- generate_density_data(unlist(param_estimates[[2]]), unlist(param_se[[2]]), x_slope)
```

```
# Plot the density for the Intercept
```

```
print(
  ggplot(density_data_intercept, aes(x, y)) +
  geom_line() +
  ggtitle("Density of Intercept Estimates") +
  xlab("Intercept Estimate") + ylab("Density")
)
```

```
# Plot the density for the Slope
```

```
print(  
  ggplot(density_data_slope, aes(x, y)) +  
  geom_line() +  
  ggtitle("Density of Slope Estimates") +  
  xlab("Slope Estimate") + ylab("Density")  
)
```

1.5 Conclusion

Conclusion: Recap your topic and provide discussion on relevant points from your presentation. Perhaps motivate an extension or different application to which your Model/Process could be utilized.

1.6 Reference

References: In your presentation you do not need to have a formal reference slide, however, be sure to appropriately reference books/papers/codes that you used in your written report.

1. Lai, Mark. 2019. "Course Handouts for Bayesian Data Analysis Class." Class handouts, December 13, 2019. https://bookdown.org/marklhc/notes_bookdown/missing-data.html
2. Little, Roderick. 2011. "Calibrated Bayes, for Statistics in General, and Missing Data in Particular1." *Statistical Science* 26 (2): 162–74. <https://doi.org/10.1214/10-STS318>.